

CS186 Discussion 6 – Joins

Tables:

Companies: (company_id, industry, ipo_date)

Nyse: (company_id, date, trade, quantity)

We have 20 pages of memory, and we want to join two tables Companies and NYSE on

C.company_id = N.company_id

Attribute company_id is the primary key for Companies.

For every tuple in Companies, assume there are 4 matching tuples in NYSE.

NYSE contains [N] = 100 pages, NYSE holds pN = 100 tuples per page.

Companies contains [C] = 50 pages, C holds pC = 50 tuples per page.

These are unclustered B+ indexes on C.company_id and N.company_id

For both indexes, assume it takes 2 I/Os to access a leaf.

1. How many disk I/Os are needed to perform a simple nested loops join?
Using Companies as the outer relation yields the lower I/O count
 $[C] + pC[C]*[N] = 50 + 50*50*100 = 50 + 50*50*100 = 250,050$ I/Os
2. Index nested loop join?
With C as the outer relation (meaning we probe on the index on N):
 $[C] + [C] * pC * (\text{cost to find matching NYSE tuples})$
 $= 50 + 50 * 50 * (2 + 4) = 15,050$ I/Os
3. How about block nested loops join?
 $(\# \text{ pages in smaller relation}) + \text{ceil}[(\# \text{ pages in smaller relation}) / (\# \text{ pages in memory} - 2)] * (\# \text{ pages in larger relation})$
 $= 50 + \text{ceil}(50/18) * 100 = 350$ I/Os
4. How about sort merge join? (Assume that both tables are sorted in 2 passes and that we join during the merge pass of sorting i.e. we are using the optimized version of sort merge join)
 $3[C] + 3[N] = 450$ I/Os
5. How about a hash join? (Assume no recursive partitioning and ignore output costs)
Partitioning phase: $2([C] + [N])$
Matching phase: $[C] + [N]$
 $3[C] + 3[N] = 450$ I/Os
6. Now assume the index on NYSE.company_id is *clustered*. What is the cost of an index nested loops join using companies as the outer relation?
 $[C] + [C]*pC * (\text{cost to find matching NYSE tuples})$
 $= 50 + 50 * 50 * (2 + \lceil 4/100 \rceil) = 7,550$ I/Os